

Sciris: Simplifying scientific software in Python

Cliff C. Kerr^{1,2}✉, Paula Sanz-Leon^{1,3}, Romesh G. Abeysuriya^{1,4}, George L. Chadderdon^{4,5}, Vlad-Ştefan Harbuz⁶, Parham Saidi⁶, James Jansson⁴, Maria del Mar Quiroga⁴, Rowan Martin-Hughes⁴, Sherrie L. Kelly⁴, Jamie A. Cohen¹, Robyn M. Stuart^{1,7}, and Anna Nachesa⁸

¹ Institute for Disease Modeling, Global Health Division, Bill & Melinda Gates Foundation, Seattle, USA
² School of Physics, University of Sydney, Sydney, Australia ³ QIMR Berghofer Medical Research Institute, Brisbane, Australia ⁴ Burnet Institute, Melbourne, Australia ⁵ CAE USA, Tampa, USA ⁶ Saffron Software, Bucharest, Romania ⁷ Department of Mathematical Sciences, University of Copenhagen, Copenhagen, Denmark ⁸ Google, Zürich, Switzerland ✉ Corresponding author

DOI: 10.xxxxxx/draft

Software

- Review ✉
- Repository ✉
- Archive ✉

Editor: Open Journals ✉

Reviewers:

- @openjournals

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

Summary

Sciris aims to streamline the development of scientific software by making it easier to perform common tasks. Sciris provides classes and functions that simplify access to frequently used low-level functionality in the core libraries of the scientific Python ecosystem (such as numpy for math and matplotlib for plotting), as well as in libraries of broader scope (such as multiprocessing for parallelization and pickle for saving and loading objects). While low-level functionality is valuable for developing robust software applications, it can divert focus from the scientific problems being solved. Some of Sciris' key features include: ensuring consistent dictionary, list, and array types (e.g., enabling users to provide inputs as either lists or arrays); enabling ordered dictionary elements to be referenced by index; simplifying datetime arithmetic by allowing date input in multiple formats, including strings; simplifying the saving and loading of files and complex objects; and simplifying the parallel execution of code. Sciris makes writing scientific code in Python faster, more pleasant, and more accessible, especially for people without extensive training in software development. With Sciris, users can achieve the same functionality with fewer lines of code, avoid reinventing the wheel, and spend less time looking up recipes on Stack Overflow. Sciris also forms the basis of ScirisWeb, an additional set of tools for building Flask-based Python webapps.

Statement of need

With the increasing availability of large volumes of data and computing resources, scientists across multiple fields of research have been able to tackle increasingly complex problems. But to harness these resources, the need for domain-specific software has become much greater. As the complexity of the questions being tackled has increased, so too has the software used to answer them, creating a steep learning curve and an increasing burden of code review ([Editorial Nature Methods, 2018](#)).

Scientific code workflows (e.g., either a full cycle in the development of a new software library, or in the execution of a one-off analysis) typically rely on multiple codebases, including but not limited to: low-level libraries, domain-specific open-source software, and self-developed and/or inherited Swiss-army-knife toolboxes (whose original developer may or may not be around to pass on undocumented wisdom). Several scientific communities have adopted collaborative, community-driven, open-source software approaches due to the significant savings in development costs and increases in code quality that they afford ([Kerr, 2019](#)), such as [astropy](#) ([Robitaille et al., 2013](#)), [fmripiprep](#) ([Esteban et al., 2019](#)), and [nextstrain](#)

(Hadfield et al., 2018). Despite this progress, a large fraction of scientific software efforts remain a solo adventure. This leads to proliferation of tools where resources are largely spent reinventing wheels of variable quality, which jeopardizes the code's minimum requirements of being "re-runnable, repeatable, reproducible, reusable, and replicable" (Benureau & Rougier, 2018).

Beyond these requirements, low-level programming abstractions may get in the way of clarifying the science. For instance, one of the reasons PyTorch has become popular in academic and research environments is its success in making models easier to write compared to TensorFlow (Lorica, 2017). The need for libraries that provide "simplifying interfaces" for research applications is reflected by the development of multiple libraries in scientific Python ecosystem that have enabled researchers focus their time and efforts on solving problems, prototyping solutions, deploying applications and educating their communities. In addition to PyTorch (simplifying/extending Tensorflow), these include seaborn (simplifying/extending Matplotlib) (Waskom, 2021), pingouin (simplifying/extending pandas), and PyVista (simplifying/extending VTK) (Sullivan & Kaszynski, 2019), among many others.

Sciris (whose name comes from a combination of "scientific" and "iris", the Greek word for "rainbow") began in 2014, to support development of the Optima suite of models (Kerr et al., 2015). We repeatedly encountering the same inconveniences while building scientific webapps, and so began collecting the tools we used to overcome them into a shared library. While Python is considered an easy-to-use language for beginners, the motivation that shaped Sciris' evolution was to further lower the barriers to access the numerous supporting libraries we were using.

Thus, we believe use of Sciris will result in more effective scientific code production for solo developers and teams alike, including increased longevity (Perkel, 2020) of new scientific libraries. Some of the key functional aspects that Sciris provides are: (i) brevity through simple interfaces; (ii) "dejargonification"; (iii) fine-grained exception handling; and (iv) version management. We expand on each of these below, but first provide a vignette that illustrates many of Sciris' features.

Vignette

Compared to a domain-specific language like MATLAB, even relatively simple scientific codes in Python can require significant "boilerplate". This extra code can obscure the key logic of the scientific problem.

For example, imagine that we wish to sample random numbers from a user-defined function with varying noise levels, save the intermediate calculations, and plot the results. In vanilla Python, each of these operations is somewhat cumbersome. Figure 1 presents two functionally identical scripts; the one written with Sciris is considerably more readable and succinct.

This vignette illustrates many of Sciris' most-used features, including timing, parallelization, high-performance containers, file saving and loading, and plotting. For the lines of the script that differ, Sciris reduces the number of lines of code required from 33 to 7, a 79% decrease.

<pre> 1 # Define random wave generator 2 import numpy as np 3 4 def randwave(std, xmin=0, xmax=10, npts=50): 5 np.random.seed(int(100*std)) # Ensure differences between runs 6 a = np.cos(np.linspace(xmin, xmax, npts)) 7 b = np.random.randn(npts) 8 return a + b*std 9 10 # Other imports 11 - import time 12 - import multiprocessing as mp 13 - import pickle 14 - import gzip 15 - import matplotlib.pyplot as plt 16 - from mpl_toolkits.mplot3d import Axes3D # Unused but must be imported 17 18 # Start timing 19 - start = time.time() 20 21 # Calculate output in parallel 22 - multipool = mp.Pool(processes=mp.cpu_count()) 23 - waves = multipool.map(randwave, np.linspace(0, 1, 11)) 24 - multipool.close() 25 - multipool.join() 26 27 # Save to files 28 - filenames = [] 29 - for i, wave in enumerate(waves): 30 - filename = f'wave{i}.obj' 31 - with gzip.GzipFile(filename, 'wb') as fileobj: 32 - fileobj.write(pickle.dumps(wave)) 33 - filenames.append(filename) 34 35 # Create dict from files 36 - data_dict = {} 37 - for fname in filenames: 38 - with gzip.GzipFile(fname) as fileobj: 39 - filestring = fileobj.read() 40 - data_dict[fname] = pickle.loads(filestring) 41 42 # Create 3D plot 43 - data = np.array([data_dict[fname] for fname in filenames]) 44 - fig = plt.figure() 45 - ax = plt.axes(projection='3d') 46 - ny, nx = np.array(data).shape 47 - x = np.arange(nx) 48 - y = np.arange(ny) 49 - X, Y = np.meshgrid(x, y) 50 - surf = ax.plot_surface(X, Y, data, cmap='coolwarm') 51 - fig.colorbar(surf) 52 53 # Print elapsed time 54 - elapsed = time.time() - start 55 - print(f'Elapsed time: {elapsed:0.1f} s')</pre>	<pre> 1 # Define random wave generator 2 import numpy as np 3 4 def randwave(std, xmin=0, xmax=10, npts=50): 5 np.random.seed(int(100*std)) # Ensure differences between runs 6 a = np.cos(np.linspace(xmin, xmax, npts)) 7 b = np.random.randn(npts) 8 return a + b*std 9 10 # Other imports 11 + import sciris as sc 12 13 # Start timing 14 + t = sc.timer() 15 16 # Calculate output in parallel 17 + waves = sc.parallelize(randwave, np.linspace(0, 1, 11)) 18 19 # Save to files 20 + filenames = [sc.save(f'wave{i}.obj', wave) for i, wave in enumerate(waves)] 21 22 # Create dict from files 23 + data = sc.odict({fname:sc.load(fname) for fname in filenames}) 24 25 # Create 3D plot 26 + sc.surf3d(data[:,], plotkwargs=dict(cmap='orangeblue')) 27 28 # Print elapsed time 29 + t.toc()</pre>
---	---

Figure 1: Comparison of a functionally identical script without Sciris (left) and with Sciris (right), showing a nearly five-fold reduction in lines of code. The resulting plot is shown in Figure 2.

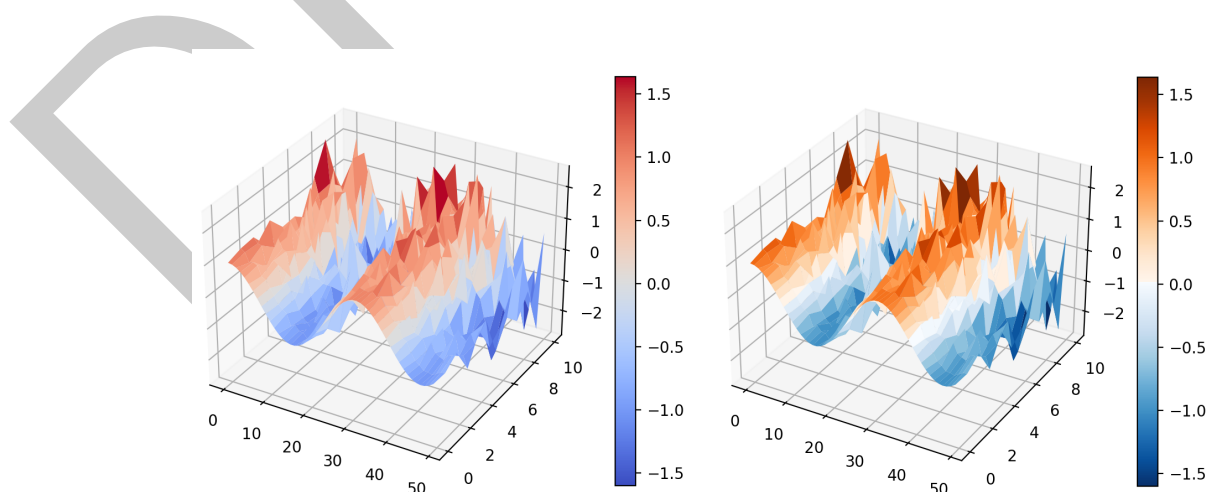


Figure 2: Output of the codes shown in Figure 1, without Sciris (left) and with Sciris (right). The two are identical except for the new high-contrast colormap available in Sciris.

Design philosophy

The aim of Sciris is to make common tasks simple. Sciris includes implementations of heavily used code patterns and abstractions that facilitate the development and deployment of complex domain-specific scientific applications, and helps non-specialist audiences interact with these applications. We note that Sciris “stands on the shoulders of giants”, and as such is not intended as a replacement of these libraries, but rather as an interface that facilitate a more effective and sustainable development process through the following principles:

Brevity through simple interfaces. Sciris packages common patterns requiring multiple lines of code into single, simple functions. With these functions one can succinctly express and execute frequent plotting tasks (e.g., `sc.commaticks`, `sc.dateformatter`, `sc.plot3d`); ensure consistent types and merging containers (e.g., `sc.toarray`, `sc.mergedicts`, `sc.mergelists`), or even perform line-by-line performance profiling (`sc.profile`). Brevity is also achieved by extending functionality of well established objects (e.g., `OrderedDict` via `sc.odict`) or methods (e.g., `isinstance` via `sc.checktype` that enables the comparison of objects against higher-level types like `arraylike`).

Dejargonification. Sciris aims to use plain function names (e.g., `sc.smooth`, `sc.findnearest`, `sc.safedivide`) so that the resulting code is as scientifically clear and human-readable as possible. Sciris also provides some **MATLAB**-like functionality, and uses the same names (e.g., `sc.tic` and `sc.toc`; `sc.boxoff`) to minimize the learning curve for scientists with this background.

Fine-grained exception handling. Across many classes and functions, Sciris uses the keyword `die`, enabling users to set a locally scoped level of strictness in the handling of exceptions. If `die=False`, Sciris is more forgiving and softly handles exceptions by using its default (opinionated) behavior, such as printing a warning and returning `None` so users can decide how to proceed. If `die=True`, it directly raises the corresponding exception and message. This flexibility reduces the need for try-catch blocks, which can distract from the code’s scientific logic.

Version management. Keeping track of dates, authors, and code versions, plus additional notes or comments, is an essential part of scientific projects. Sciris provides methods to easily save and load metadata to/from figure files, including Git information (`sc.savefig`, `sc.gitinfo`, `sc.loadmetadata`), as well as shortcuts for comparing module versions (`sc.compareversions`) or requiring them (`sc.require`).

Sciris in practice

Our investments in Sciris paid off when in early 2020 its combination of brevity and simplicity proved crucial in enabling the rapid development of the Covasim model of COVID-19 transmission (Kerr et al., 2021). Covasim’s relative simplicity and readability, based in large part on its heavy use of Sciris, enabled it to become one of the most widely adopted models of COVID-19, used by students, researchers and policymakers in over 30 countries (Kerr et al., 2022).

In addition to Covasim, Sciris is currently used by a number of other scientific software tools, such as **Optima HIV** (Kerr et al., 2015), **Optima Nutrition** (Pearson et al., 2018), the **Cascade Analysis Tool** (Kedziora et al., 2019), **Atomica** (The Atomica Team, 2020), **Optima TB** (Goscé et al., 2021), the **Health Interventions Prioritization Tool** (Fraser-Hurt et al., 2021), **SynthPops** (Mistry et al., in preparation), and **FPsim** (O’Brien et al., 2022).

Examples of key features

Here we illustrate a smattering of key features in greater detail; further information can be found at docs.sciris.org. Figure 3 illustrates the functional modules of Sciris. Sciris is available

128 on pip (pip install sciris).

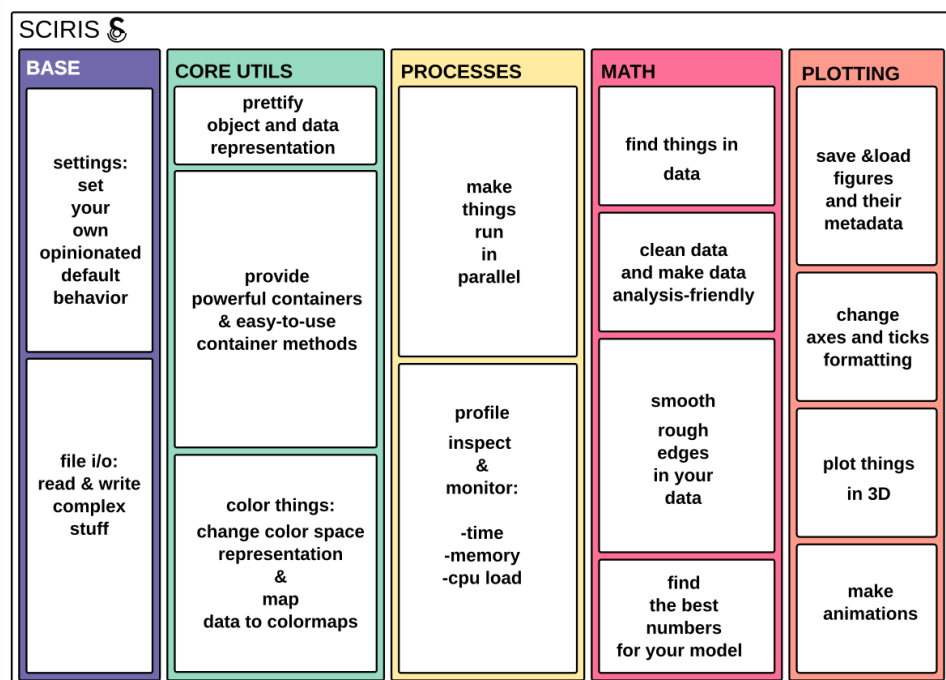


Figure 3: Block diagram of the Sciris' functionality, grouped by high-level concepts and types of tasks that are commonly performed in scientific code.

129 High-performance containers

130 One of the key features in Sciris is `sc.odict`, a flexible container representing an associative
131 array with the best-of-all-worlds features of lists, dictionaries, and numeric arrays. This is
132 based on `OrderedDict` from [collections](#), but supports list methods like integer indexing, key
133 slicing, and item insertion:

```
> data = sc.odict(a=[1,2,3], b=[4,5,6])
> data['a'] == data[0]
> data[:].sum() == 21
> for i, key, value in data.enumitems():
    print(f'Item {i} is named "{key}" and has value {value}')
Item 0 is named "a" and has value [1, 2, 3]
Item 1 is named "b" and has value [4, 5, 6]
```

134 Numerical utilities

135 Indexing arrays is a common task in NumPy, but can be difficult due to incompatibilities of
136 object type. `sc.findinds` will find matches even if two things are not exactly equal due to
137 differences in numeric type (e.g., floats vs. integers, lists vs. arrays). The code shown below
138 produces the same result as calling `np.nonzero(np.isclose(arr, val))[0]`.

```
> sc.findinds([2,3,6,3], 3.0)
array([1,3])
```

139 Parallelization

140 A frequent hurdle scientists face is parallelization. Sciris provides `sc.parallelize`, which acts
141 as a shortcut for using `multiprocessing.Pool()`. By default it adjusts the pool size based on the

142 CPUs available, but can also use either a fixed number of CPUs or allocate them dynamically
143 based on load (`sc.loadbalancer`). This example shows three equivalent ways to iterate over
144 multiple complex arguments:

```
> def f(x, y):
>     return x*y

> out1 = sc.parallelize(func=f, iterarg=[(1,2),(2,3),(3,4)])
> out2 = sc.parallelize(func=f, iterkwargs={'x':[1,2,3], 'y':[2,3,4]})
> out3 = sc.parallelize(func=f, iterkwargs=[{'x':1, 'y':2},
                                             {'x':2, 'y':3},
                                             {'x':3, 'y':4}])
```

145 Plotting

146 Numerous shortcuts for customizing and prettifying plots are available in Sciris; several
147 commonly used features are illustrated below, with the results shown in Figure 4:

```
> sc.options(font='Garamond') # Set custom font
> x = sc.daterange('2022-06-01', '2022-12-31', as_date=True) # Create dates
> y = sc.smooth(np.random.randn(len(x))*2)*1000 # Create smoothed random numbers
> c = sc.vectocolor(y, cmap='turbo') # Set colors proportional to y values
> plt.scatter(x, y, c=c) # Plot the data
> sc.dateformatter() # Automatic x-axis date formatter
> sc.commaticks() # Add commas to y-axis tick labels
> sc.setylim() # Automatically set the y-axis limits
> sc.boxoff() # Remove the top and right axis spines
```




Figure 4: Example of plot customizations via Sciris, including x- and y-axis tick labels and the font.

148 ScirisWeb

149 While a full description of [ScirisWeb](#) is beyond the scope of this paper, briefly, it builds on
150 Sciris to enable the rapid development of Python-based webapps, including those powering
151 [Covasim](#) and [Optima Nutrition](#). By default, ScirisWeb uses [Vuejs](#) and [sciris-js](#) for the frontend,
152 [Flask](#) as the web framework, [Redis](#) for the (optional) database, and Matplotlib/[mpld3](#) for
153 plotting. However, ScirisWeb is completely modular, which means that it could also be used to

(for example) link a [React](#) frontend to a [MySQL](#) database with [Plotly](#) figures. This modularity is in contrast to full-stack solutions such as [Plotly Dash](#), [Streamlit](#), and [Voilà](#). While these libraries are even easier to use than ScirisWeb (since they do not require any knowledge of JavaScript), they provide limited options for customization or switching between technology stacks. In contrast, ScirisWeb provides the flexibility of a custom-written webapp within the context of an “it just works” framework.

Acknowledgements

The Sciris Development Team (info@sciris.org) wishes to thank David J. Kedziora, Dominic Delpont, Kevin M. Jablonka, Meikang Wu, and Dina Mistry for providing helpful feedback on the Sciris library. David P. Wilson, William B. Lytton, and Daniel J. Klein provided in-kind support of Sciris development. Financial and personnel support has been provided by the United States [Defense Advanced Research Projects Agency](#) (DARPA) Contract N66001-10-C-2008 (2010–2014), World Bank Assignment 1045478 (2011–2015), the Australian National Health and Medical Research Council (NHMRC) Project Grant APP1086540 (2015–2017), the Australian Research Council (ARC) Discovery Early Career Research Award (DECRA) Fellowship Grant DE140101375 (2014–2019), Intellectual Ventures (2019–2020), and the Bill & Melinda Gates Foundation (2020–present).

References

- Benureau, F. C., & Rougier, N. P. (2018). Re-run, repeat, reproduce, reuse, replicate: Transforming code into scientific contributions. *Frontiers in Neuroinformatics*, 11, 69.
- Editorial Nature Methods. (2018). Easing the burden of code review. *Nature Methods*, 15(9), page641.
- Esteban, O., Markiewicz, C. J., Blair, R. W., Moodie, C. A., Isik, A. I., Erramuzpe, A., Kent, J. D., Goncalves, M., DuPre, E., Snyder, M., & others. (2019). fMRIPrep: A robust preprocessing pipeline for functional MRI. *Nature Methods*, 16(1), 111–116.
- Fraser-Hurt, N., Hou, X., Wilkinson, T., Duran, D., Abou Jaoude, G. J., Skordis, J., Chukwuma, A., Lao Pena, C., Tshivuila Matala, O. O., Gorgens, M., & others. (2021). Using allocative efficiency analysis to inform health benefits package design for progressing towards universal health coverage: Proof-of-concept studies in countries seeking decision support. *PLOS One*, 16(11), e0260247.
- Goscé, L., Abou Jaoude, G. J., Kedziora, D. J., Benedikt, C., Hussain, A., Jarvis, S., Skrahina, A., Klimuk, D., Hurevich, H., Zhao, F., & others. (2021). Optima TB: A tool to help optimally allocate tuberculosis spending. *PLOS Computational Biology*, 17(9), e1009255.
- Hadfield, J., Megill, C., Bell, S. M., Huddleston, J., Potter, B., Callender, C., Sagulenko, P., Bedford, T., & Neher, R. A. (2018). Nextstrain: Real-time tracking of pathogen evolution. *Bioinformatics*, 34(23), 4121–4123.
- Kedziora, D. J., Abeyesuriya, R., Kerr, C. C., Chadderdon, G. L., Harbuz, V.-Ş., Metzger, S., Wilson, D. P., & Stuart, R. M. (2019). The cascade analysis tool: Software to analyze and optimize care cascades. *Gates Open Research*, 3.
- Kerr, C. C. (2019). Is epidemiology ready for big software? *Pathogens and Disease*, 77(1), ftz006.
- Kerr, C. C., Stuart, R. M., Gray, R. T., Shattock, A. J., Fraser-Hurt, N., Benedikt, C., Haacker, M., Berdnikov, M., Mahmood, A. M., Jaber, S. A., & others. (2015). Optima: A model for HIV epidemic analysis, program prioritization, and resource optimization. *Journal of Acquired Immune Deficiency Syndromes*, 69(3), 365–376.

- 199 Kerr, C. C., Stuart, R. M., Mistry, D., Abeysuriya, R. G., Cohen, J. A., George, L., Jastrzebski,
200 M., Famulare, M., Wenger, E., & Klein, D. J. (2022). Python vs. The pandemic: A case
201 study in high-stakes software development. *Proceedings of the 21st Python in Science*
202 *Conference (SciPy 2022)*.
- 203 Kerr, C. C., Stuart, R. M., Mistry, D., Abeysuriya, R. G., Rosenfeld, K., Hart, G. R., Núñez, R.
204 C., Cohen, J. A., Selvaraj, P., Hagedorn, B., & others. (2021). Covasim: An agent-based
205 model of COVID-19 dynamics and interventions. *PLOS Computational Biology*, 17(7),
206 e1009149.
- 207 Lorica, B. (2017). *Why AI and machine learning researchers are beginning to embrace Py-*
208 *Torch*. [oreilly.com/radar/podcast/why-ai-and-machine-learning-researchers-are-beginning-](https://oreilly.com/radar/podcast/why-ai-and-machine-learning-researchers-are-beginning-to-embrace-pytorch)
209 [to-embrace-pytorch](https://oreilly.com/radar/podcast/why-ai-and-machine-learning-researchers-are-beginning-to-embrace-pytorch)
- 210 Mistry, D., Kerr, C. C., Abeysuriya, R. G., Wu, M., Fisher, M., Thompson, A., Skrip, L.,
211 Cohen, J. A., & Klein, D. J. (in preparation). *SynthPops: A generative model of human*
212 *contact networks*.
- 213 O'Brien, M. L., Valente, A., Chabot-Couture, G., Proctor, J., Klein, D., Kerr, C., & Zimmer-
214 mann, M. (2022). FPSim: An agent-based model of family planning for informed policy
215 decision-making. *PAA 2022 Annual Meeting*.
- 216 Pearson, R., Killedar, M., Petravic, J., Kakietek, J. J., Scott, N., Grantham, K. L., Stuart,
217 R. M., Kedziora, D. J., Kerr, C. C., Skordis-Worrall, J., & others. (2018). Optima
218 nutrition: An allocative efficiency tool to reduce childhood stunting by better targeting of
219 nutrition-related interventions. *BMC Public Health*, 18(1), 1–12.
- 220 Perkel, J. M. (2020). Challenge to scientists: Does your ten-year-old code still run? *Nature*,
221 584(7822), 656–659.
- 222 Robitaille, T. P., Tollerud, E. J., Greenfield, P., Droettboom, M., Bray, E., Aldcroft, T., Davis,
223 M., Ginsburg, A., Price-Whelan, A. M., Kerzendorf, W. E., & others. (2013). Astropy: A
224 community python package for astronomy. *Astronomy & Astrophysics*, 558, A33.
- 225 Sullivan, C., & Kaszynski, A. (2019). PyVista: 3D plotting and mesh analysis through a
226 streamlined interface for the visualization toolkit (VTK). *Journal of Open Source Software*,
227 4(37), 1450.
- 228 The Atomica Team. (2020). Atomica: A simulation engine for compartmental models. In
229 *GitHub repository*. GitHub. <https://github.com/atomicateam/atomica>
- 230 Waskom, M. L. (2021). Seaborn: Statistical data visualization. *Journal of Open Source*
231 *Software*, 6(60), 3021.